

```
>>> c++ deep-learning.cpp
```

```
>>> ./a.out data.csv
```

```
>>> 42
```

# Strategy



## OSI Model of Machine Learning

# Task Layer

$$f \left( \text{Teapot} \right) = \text{Teapot}$$

**Classification\***

$$f(\text{文内容}) = \begin{matrix} \circ \\ \vdots \\ \circ \\ \circ \\ \vdots \\ \circ \end{matrix}$$

**Translation\***

$$f \left( \text{Clouds} \right) =$$

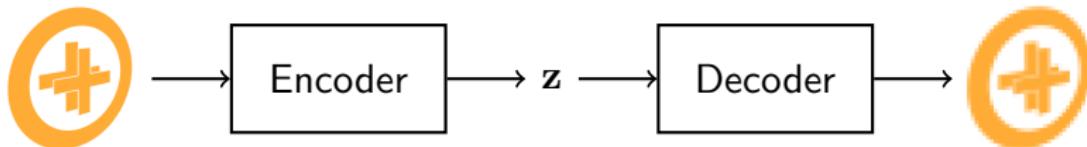


**Generative<sup>†</sup>**

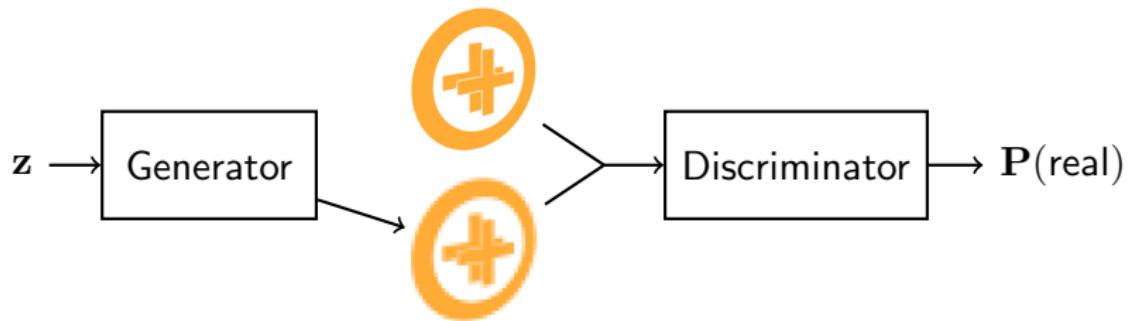
\* Supervised

† Unsupervised

## Model Layer



**Autoencoder**



**Generative Adversarial Networks (GAN)**

# Generative Adversarial Networks

$$\underbrace{\mathbf{z} \sim P_{\text{noise}}(\mathbf{z})}_{\text{Noise}} \quad \underbrace{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})}_{\text{Real Images}}$$

$$\underbrace{G(\mathbf{z})}_{\text{Fake Images}} \quad \underbrace{D(\mathbf{x})}_{\text{Probability of Being Real}}$$

# Generative Adversarial Networks

$$\underbrace{\mathbf{z} \sim P_{\text{noise}}(\mathbf{z})}_{\text{Noise}} \quad \underbrace{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})}_{\text{Real Images}}$$

$$\underbrace{G(\mathbf{z})}_{\text{Fake Images}} \quad \underbrace{D(\mathbf{x})}_{\text{Probability of Being Real}}$$

$$\frac{\int_{-\infty}^{\infty} \hat{f}(\xi) e^{2\pi i \xi x} d\xi + \sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} + \left\{ \frac{e^2}{\pi} \right\}}{\frac{1}{\sqrt{2\pi\hbar}} \int_{\text{all space}} e^{-i\mathbf{p}\cdot\mathbf{r}/\hbar} \Psi(\mathbf{r}, s_z, t) d^3\mathbf{r} - \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(x-\mu)^2}{2\sigma^2}}}$$

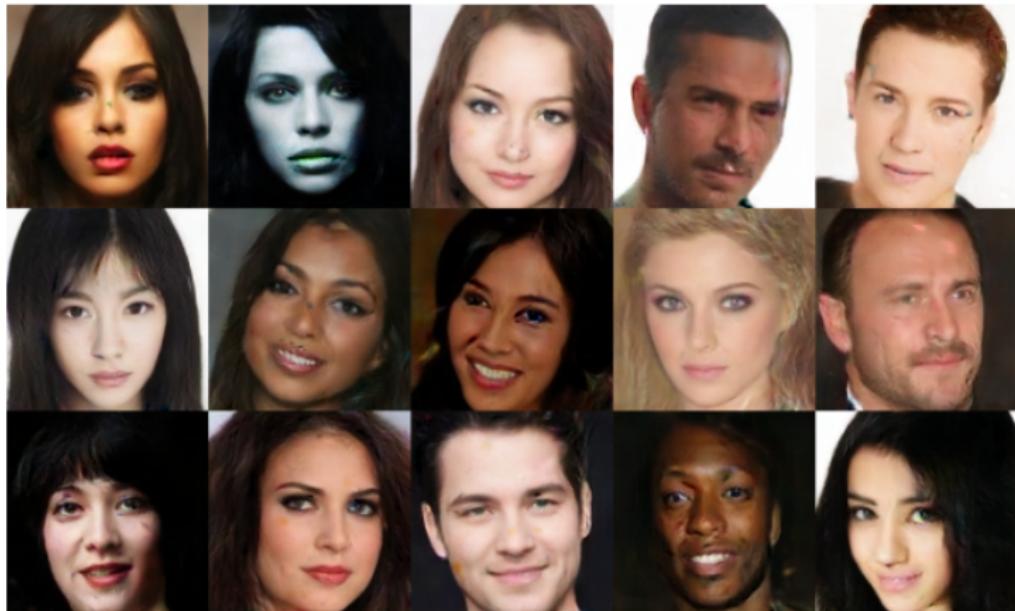
# Generative Adversarial Networks

$$\underbrace{\mathbf{z} \sim P_{\text{noise}}(\mathbf{z})}_{\text{Noise}} \quad \underbrace{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})}_{\text{Real Images}}$$

$$\underbrace{G(\mathbf{z})}_{\text{Fake Images}} \quad \underbrace{D(\mathbf{x})}_{\text{Probability of Being Real}}$$

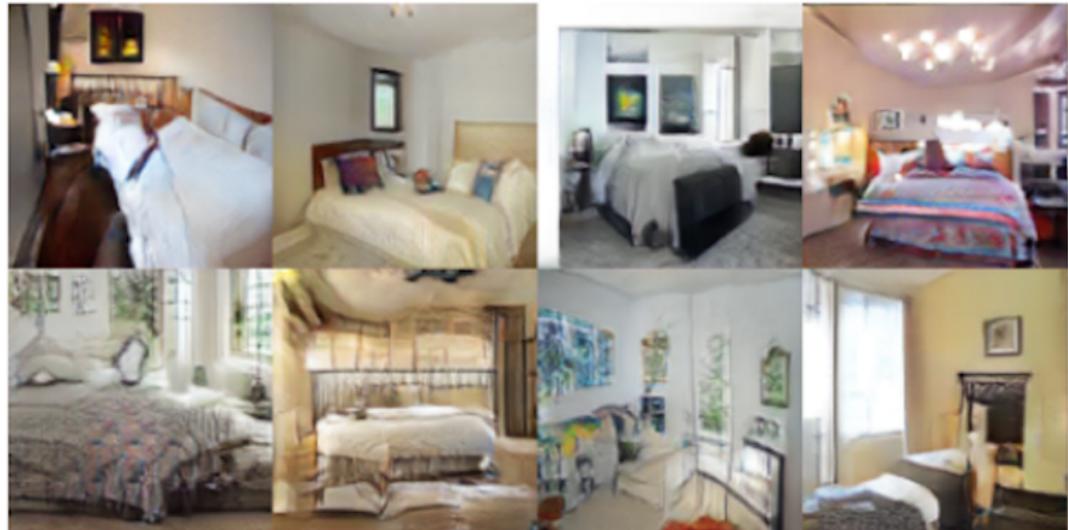
$$\min_G \max_D \log(D(\mathbf{x})) + \log(1 - D(G(\mathbf{z})))$$

# Generative Adversarial Networks



BEGAN: Berthelot et al. (2017)

# Generative Adversarial Networks



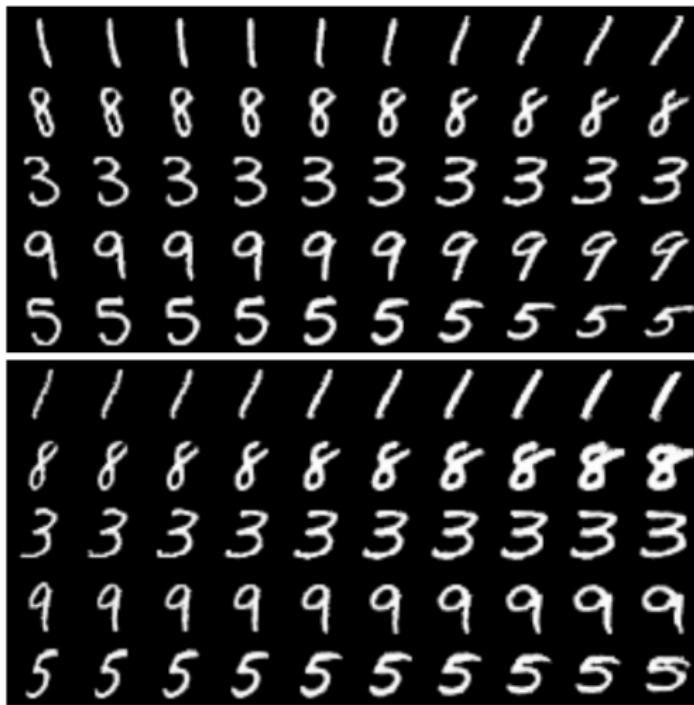
LSGAN: Mao et al. (2016)

# Generative Adversarial Networks



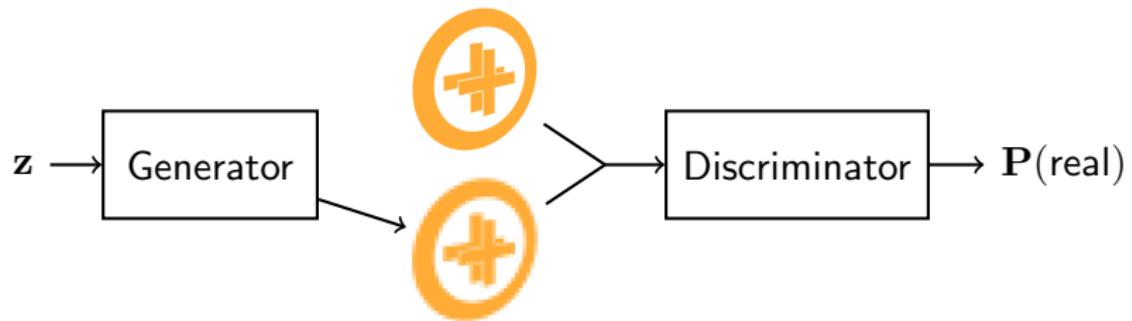
BEGAN: Berthelot et al. (2017)

# Generative Adversarial Networks



InfoGAN: Chen et al. (2016)

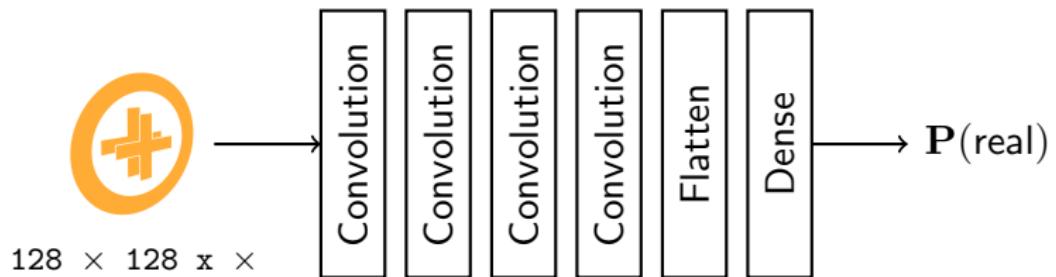
# Layer Layer



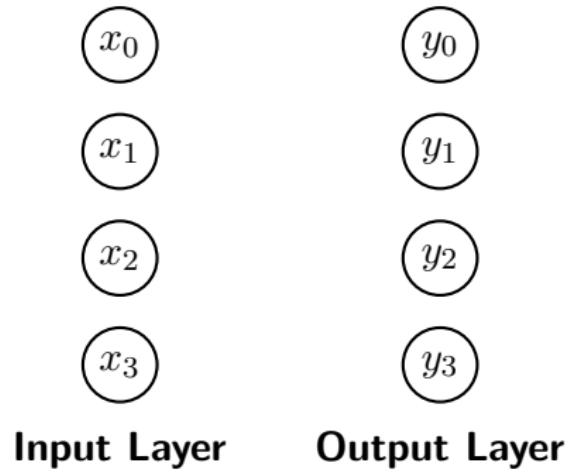
# Layer Layer



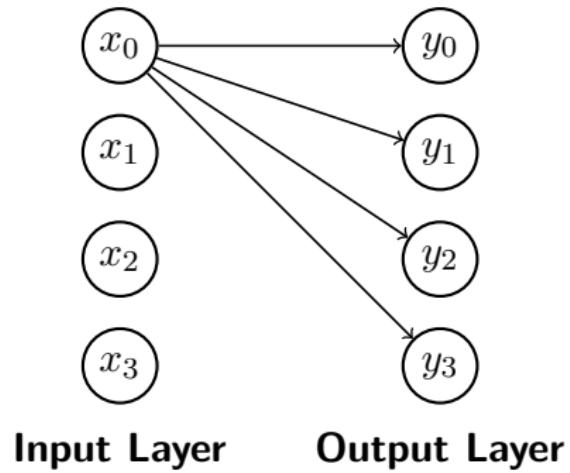
# Layer Layer



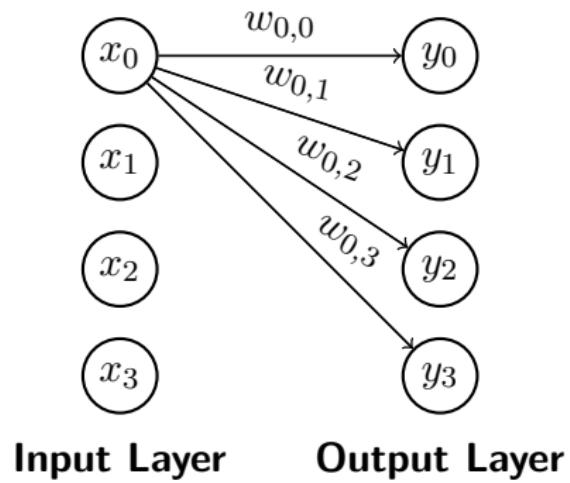
## Layers: Dense



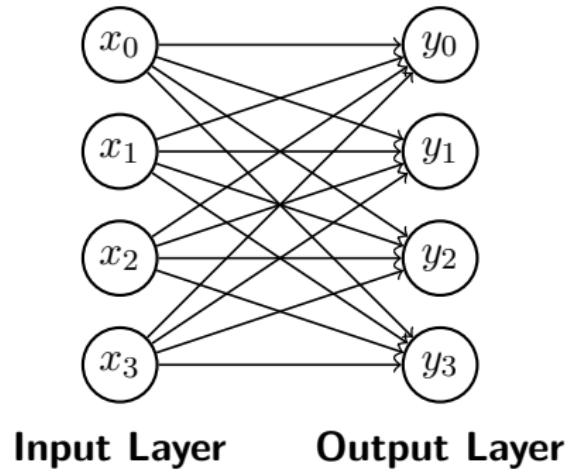
## Layers: Dense



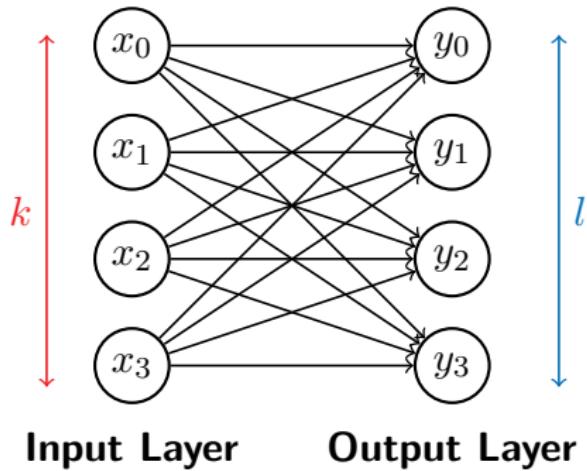
## Layers: Dense



## Layers: Dense



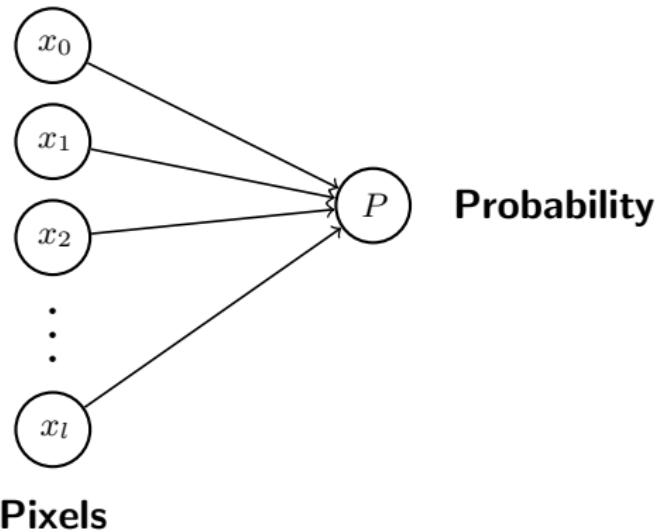
## Layers: Dense



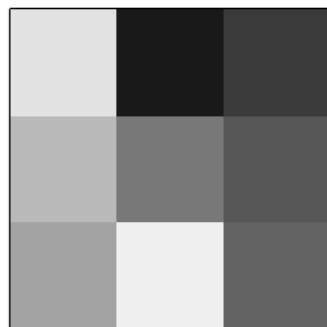
$$n \begin{bmatrix} x_{0,0} & x_{0,1} & x_{0,2} & x_{0,3} \\ x_{1,0} & x_{1,1} & x_{1,2} & x_{1,3} \\ \vdots & \vdots & \vdots & \vdots \\ x_{n,0} & x_{n,1} & x_{n,2} & x_{n,3} \end{bmatrix} \times \begin{bmatrix} w_{0,0} & w_{0,1} & w_{0,2} & w_{0,3} \\ w_{1,0} & w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,0} & w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,0} & w_{3,1} & w_{3,2} & w_{3,3} \end{bmatrix} = \begin{bmatrix} y_{0,0} & y_{0,1} & y_{0,2} & y_{0,3} \\ y_{1,0} & y_{1,1} & y_{1,2} & y_{1,3} \\ \vdots & \vdots & \vdots & \vdots \\ y_{n,0} & y_{n,1} & y_{n,2} & y_{n,3} \end{bmatrix}$$

The diagram shows a matrix multiplication operation. The input layer is represented as a column vector of size  $n$  (number of input nodes). The weight matrix has  $k$  columns (number of input nodes) and  $l$  rows (number of output nodes). The result is a column vector of size  $l$  (number of output nodes).

## Layers: Dense



# Layers: Convolution



Image

## Layers: Convolution

0.4	0.9	0.1
0.7	0.2	0.6
0.8	0.3	0.5

Image

## Layers: Convolution

0.4	0.9	0.1
0.7	0.2	0.6
0.8	0.3	0.5

Image

5.7	2.4
3.1	0.9

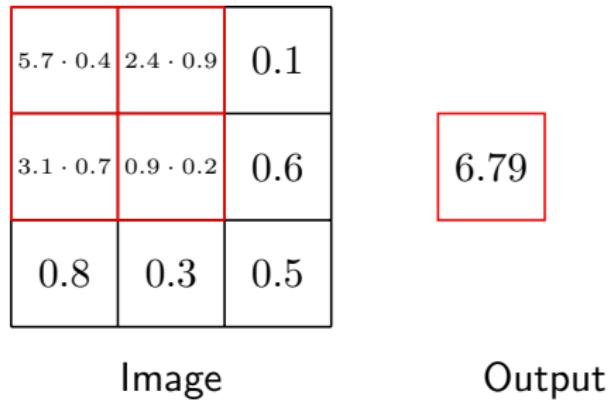
Kernel

# Layers: Convolution

$5.7 \cdot 0.4$	$2.4 \cdot 0.9$	0.1
$3.1 \cdot 0.7$	$0.9 \cdot 0.2$	0.6
0.8	0.3	0.5

Image

## Layers: Convolution



# Layers: Convolution

0.4	$5.7 \cdot 0.9$	$2.4 \cdot 0.1$
0.7	$3.1 \cdot 0.2$	$0.9 \cdot 0.6$
0.8	0.3	0.5

Image

6.79

Output

# Layers: Convolution

0.4	$5.7 \cdot 0.9$	$2.4 \cdot 0.1$
0.7	$3.1 \cdot 0.2$	$0.9 \cdot 0.6$
0.8	0.3	0.5

Image

6.79	6.53
------	------

Output

# Layers: Convolution

0.4	0.9	0.1
$5.7 \cdot 0.7$	$2.4 \cdot 0.2$	0.6
$3.1 \cdot 0.8$	$0.9 \cdot 0.3$	0.5

Image

6.79	6.53
------	------

Output

# Layers: Convolution

0.4	0.9	0.1
$5.7 \cdot 0.7$	$2.4 \cdot 0.2$	0.6
$3.1 \cdot 0.8$	$0.9 \cdot 0.3$	0.5

Image

6.79	6.53
7.67	

Output

# Layers: Convolution

0.4	0.9	0.1
0.7	$5.7 \cdot 0.2$	$2.4 \cdot 0.6$
0.8	$3.1 \cdot 0.3$	$0.9 \cdot 0.5$

Image

6.79	6.53
7.67	

Output

# Layers: Convolution

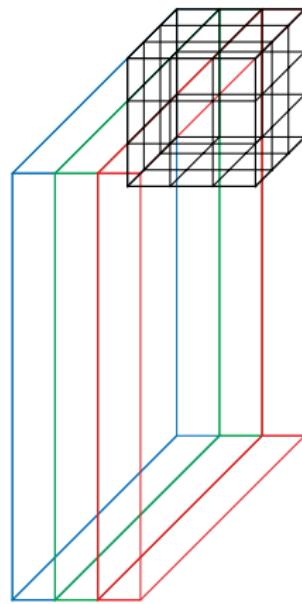
0.4	0.9	0.1
0.7	$5.7 \cdot 0.2$	$2.4 \cdot 0.6$
0.8	$3.1 \cdot 0.3$	$0.9 \cdot 0.5$

Image

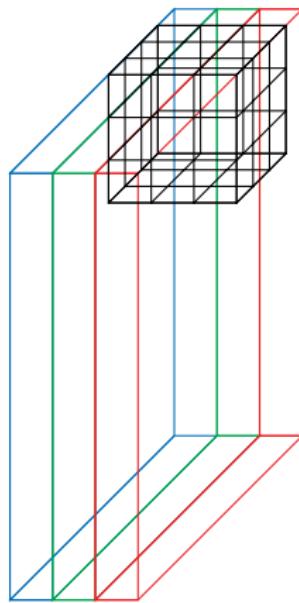
6.79	6.53
7.67	3.96

Output

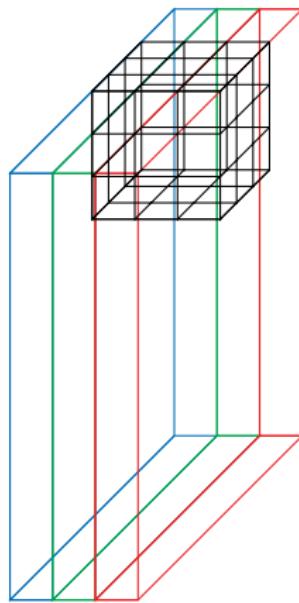
## Layers: Convolution



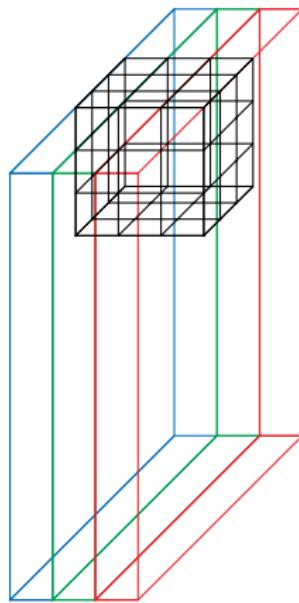
# Layers: Convolution



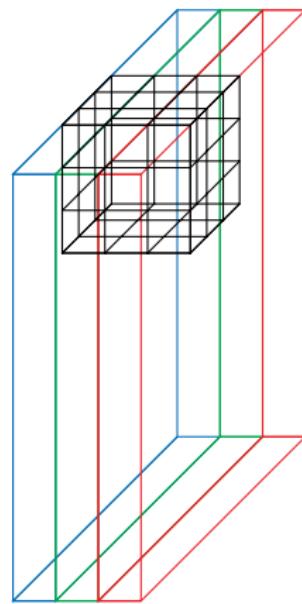
# Layers: Convolution



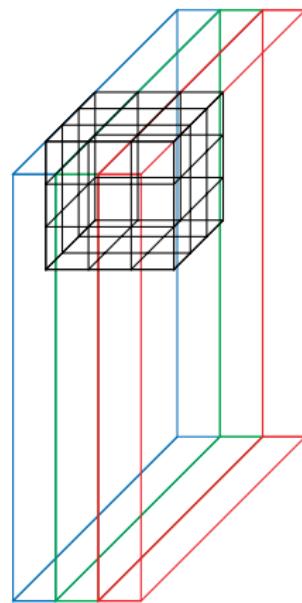
## Layers: Convolution



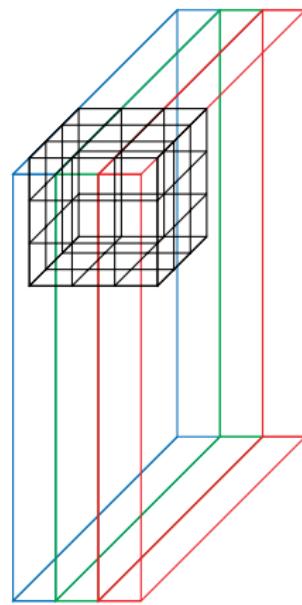
# Layers: Convolution



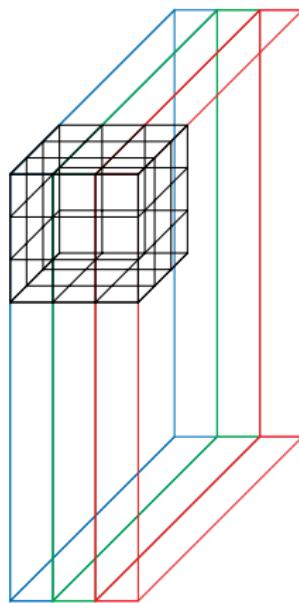
## Layers: Convolution



# Layers: Convolution



# Layers: Convolution



# Layers: Convolution

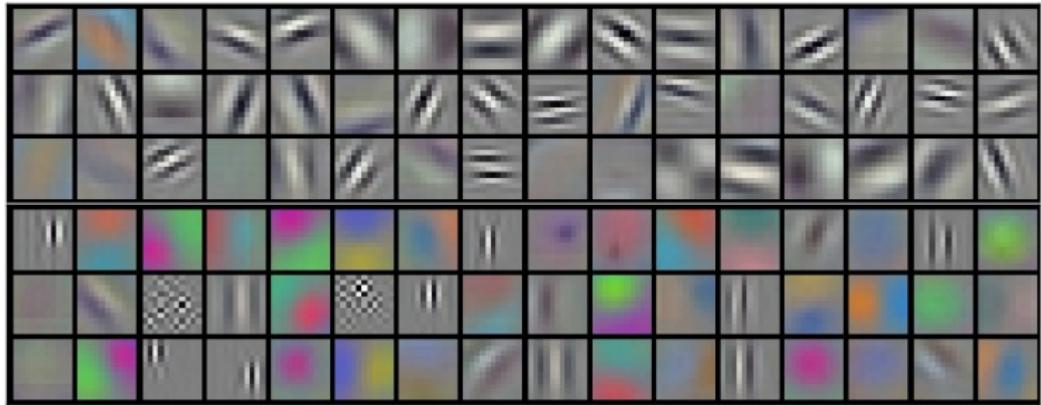
0	1	0
0	1	0
0	1	0

**Patterns**

0	0	0
-1	1	0
0	0	0

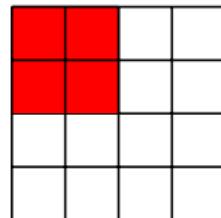
**Edges**

## Layers: Convolution

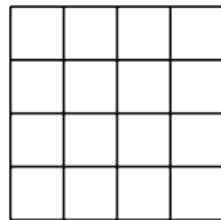


AlexNet: Krizhevsky et al. (2012)

## Layer: Convolution

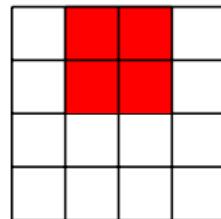


**Stride 1**

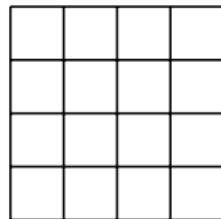


**Stride 2**

## Layer: Convolution

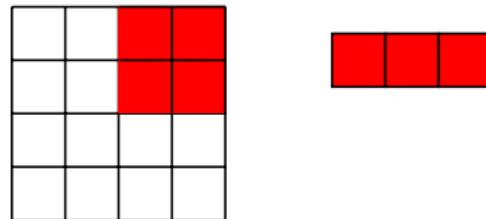


**Stride 1**

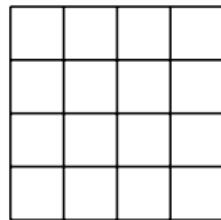


**Stride 2**

## Layer: Convolution

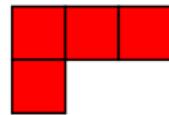
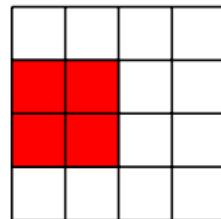


**Stride 1**

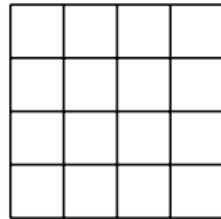


**Stride 2**

## Layer: Convolution

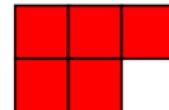
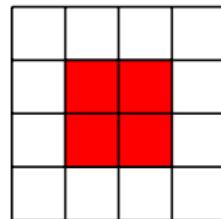


**Stride 1**

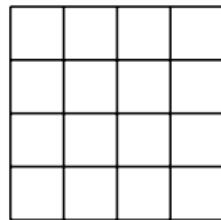


**Stride 2**

## Layer: Convolution

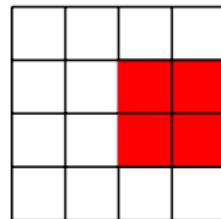


**Stride 1**

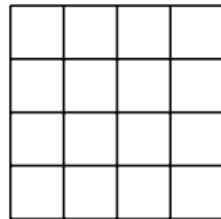


**Stride 2**

## Layer: Convolution

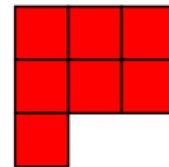
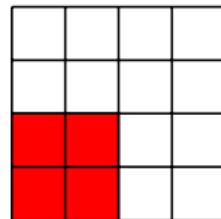


**Stride 1**

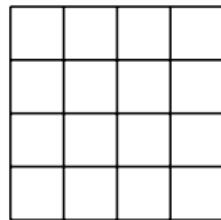


**Stride 2**

## Layer: Convolution

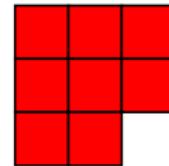
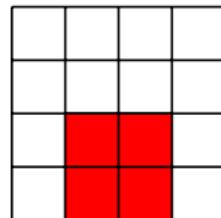


**Stride 1**

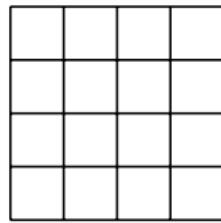


**Stride 2**

## Layer: Convolution

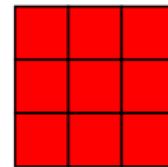
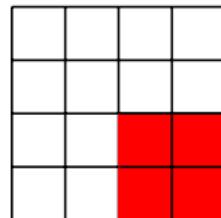


**Stride 1**

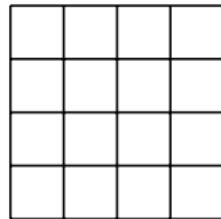


**Stride 2**

## Layer: Convolution

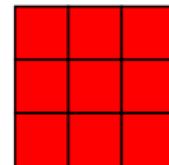
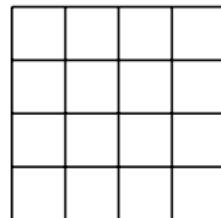


**Stride 1**

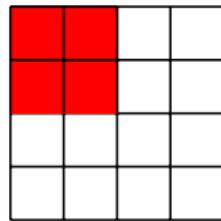


**Stride 2**

## Layer: Convolution

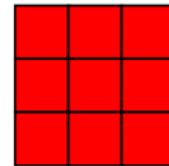
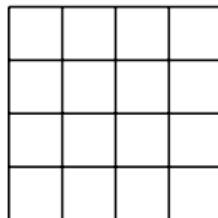


**Stride 1**

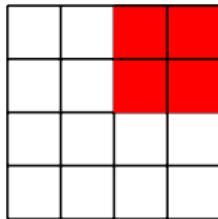


**Stride 2**

## Layer: Convolution

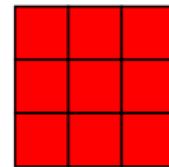
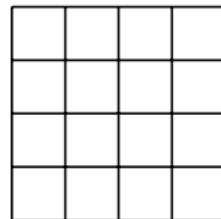


**Stride 1**

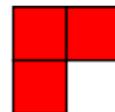
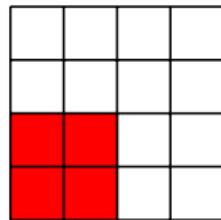


**Stride 2**

## Layer: Convolution

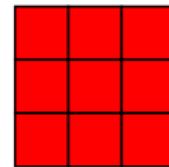
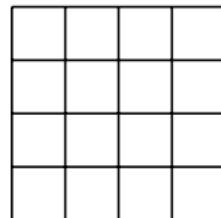


**Stride 1**

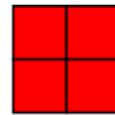
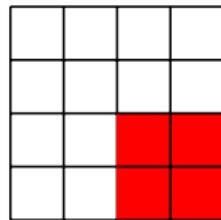


**Stride 2**

## Layer: Convolution

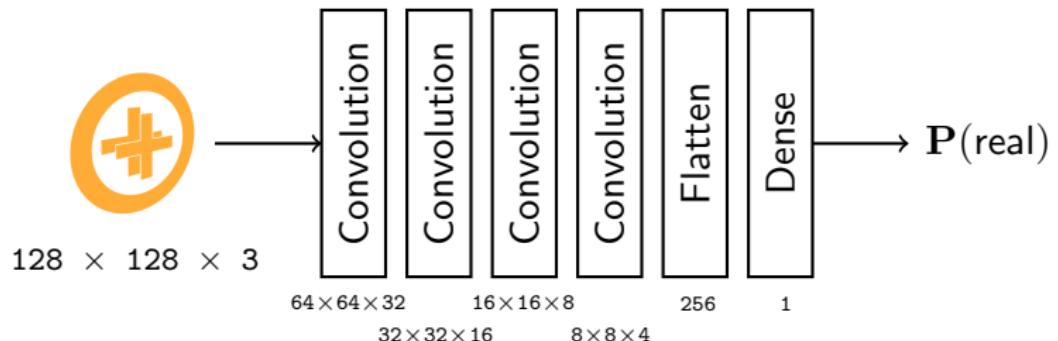


**Stride 1**



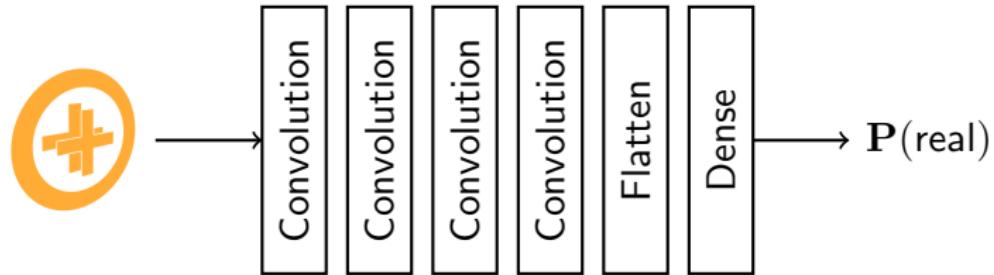
**Stride 2**

## Layer: Convolution

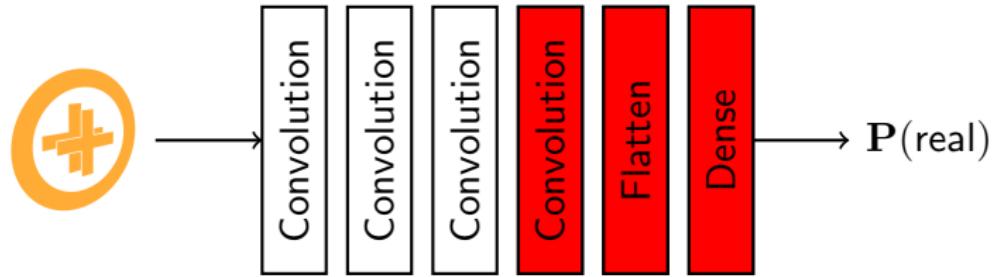


**Kernel:**  $5 \times 5$   
**Stride:** 2

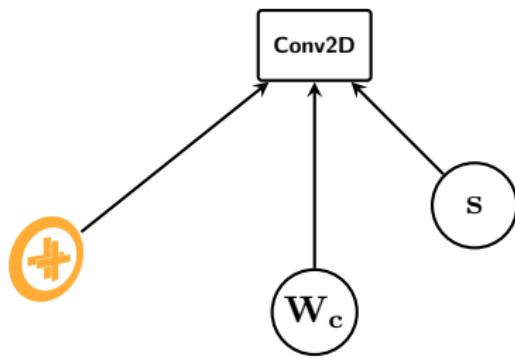
# Graph Layer



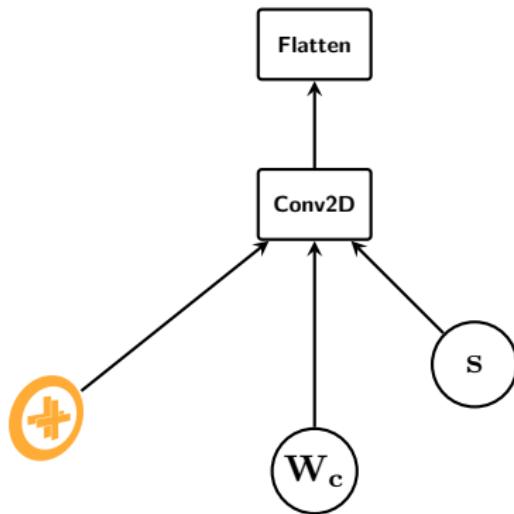
# Graph Layer



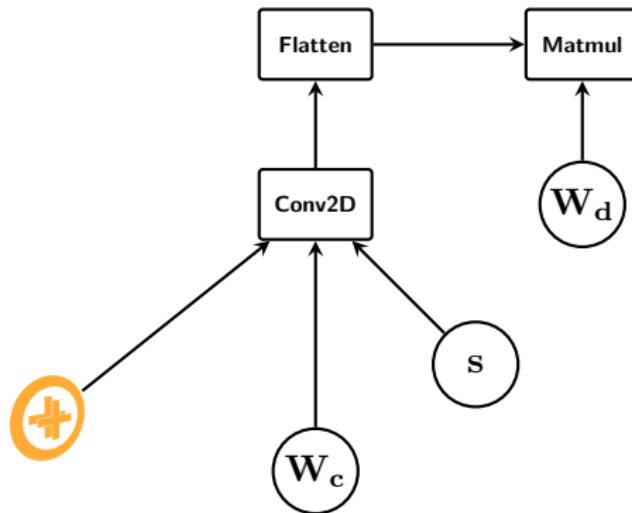
# Graph Layer



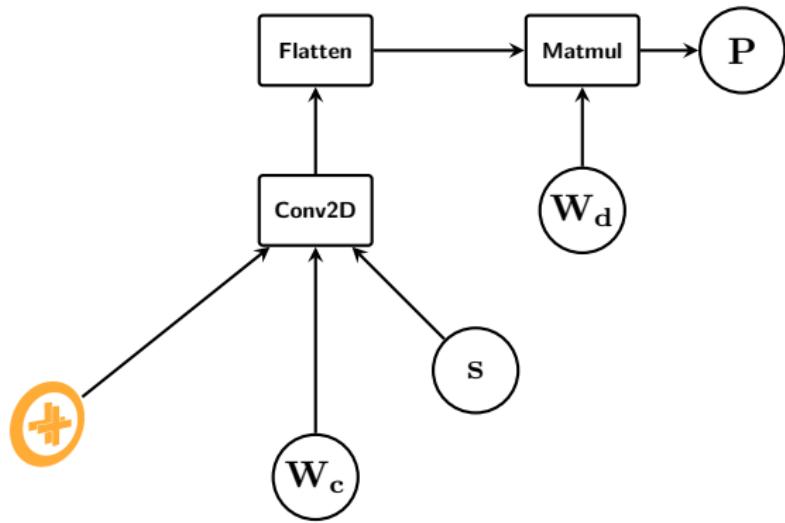
# Graph Layer



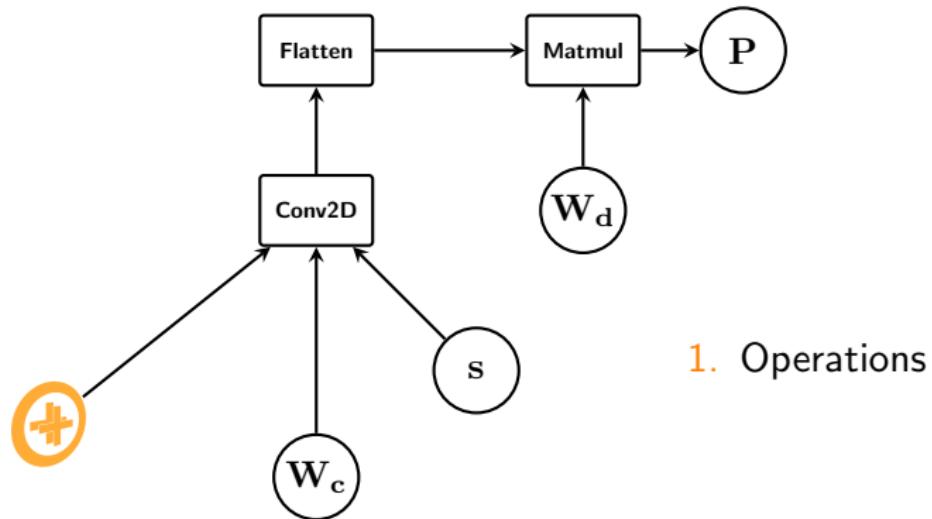
# Graph Layer



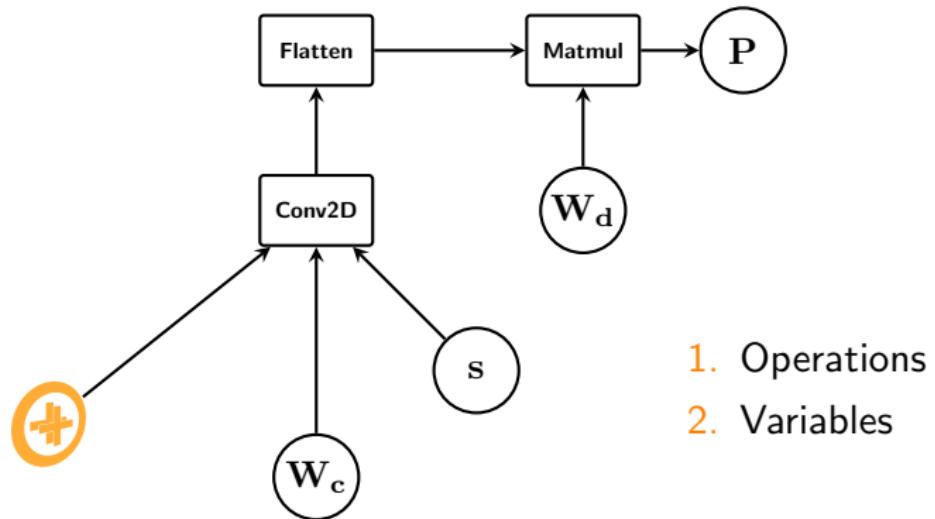
# Graph Layer



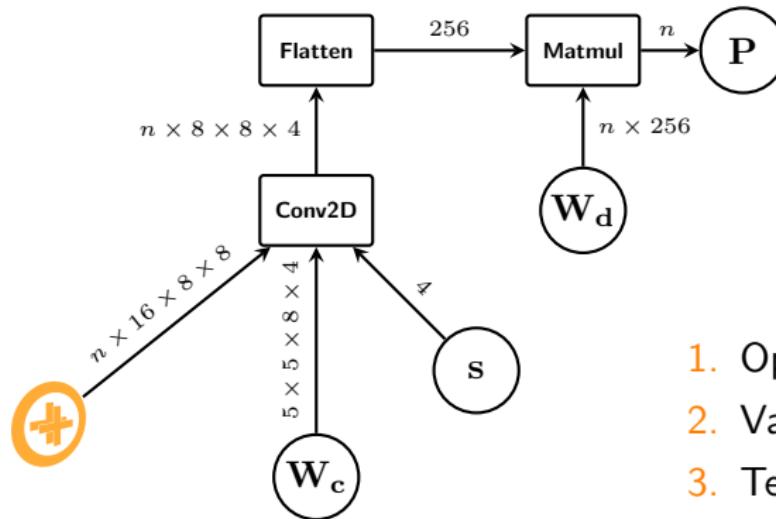
# Graph Layer



# Graph Layer

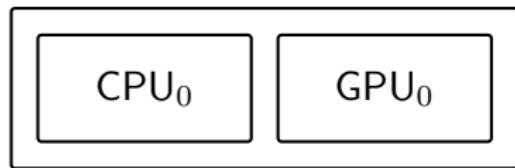
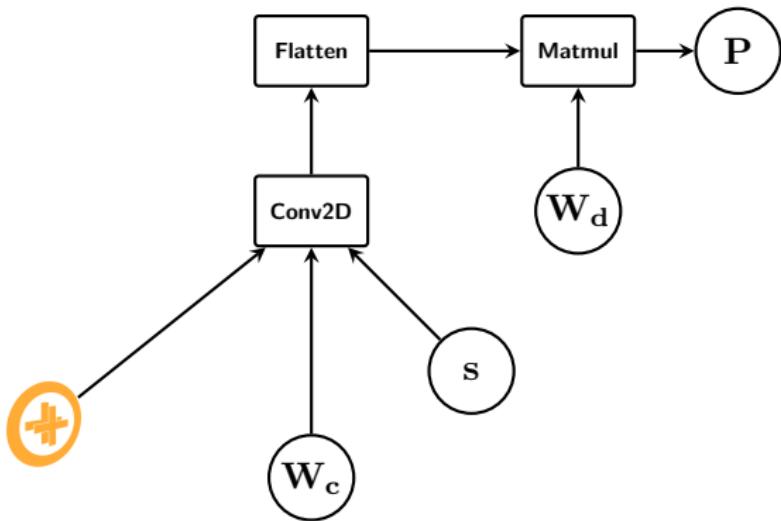


# Graph Layer

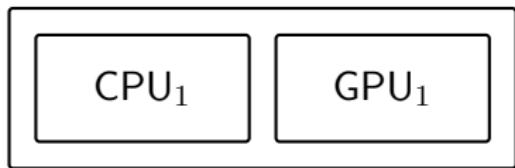


1. Operations
2. Variables
3. Tensors

## Graph Layer

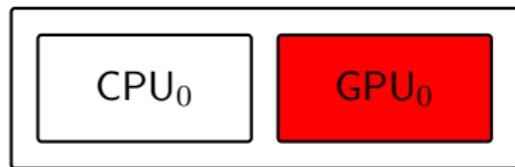
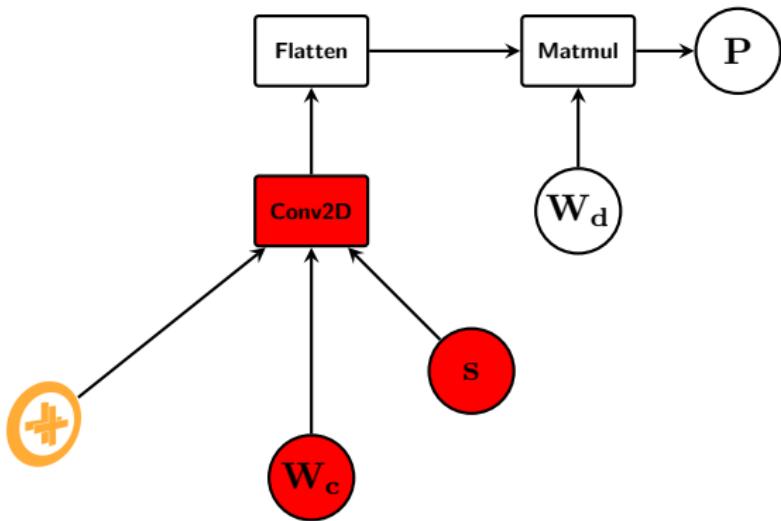


**Server<sub>0</sub>**

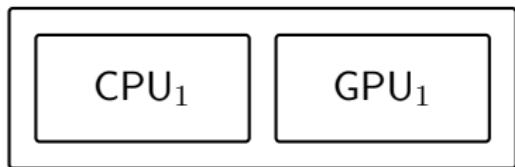


**Server<sub>1</sub>**

## Graph Layer

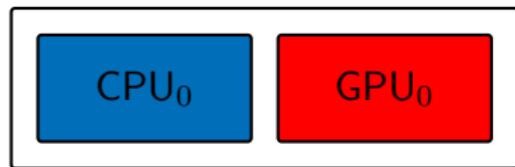
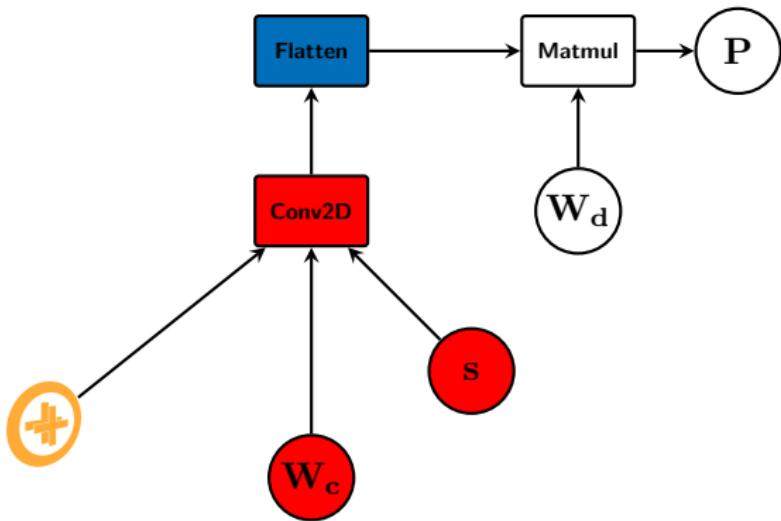


**Server<sub>0</sub>**

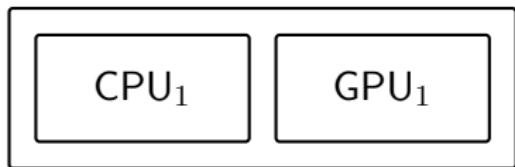


**Server<sub>1</sub>**

# Graph Layer

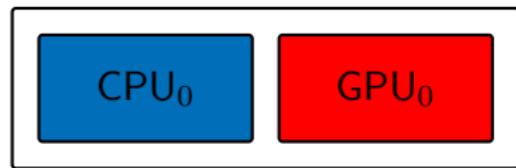
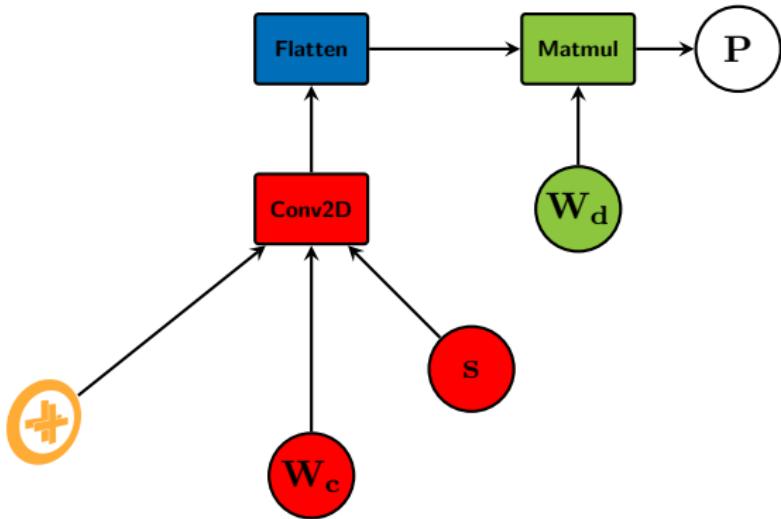


**Server<sub>0</sub>**



**Server<sub>1</sub>**

# Graph Layer

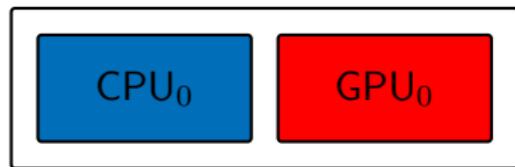
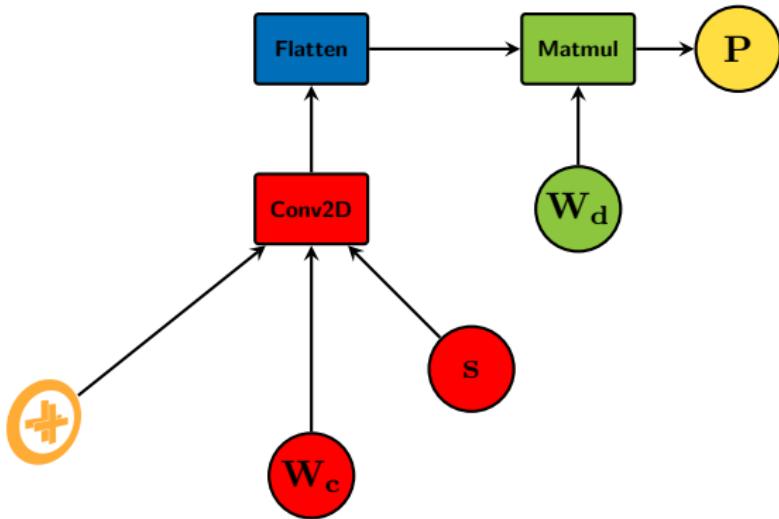


**Server<sub>0</sub>**



**Server<sub>1</sub>**

# Graph Layer

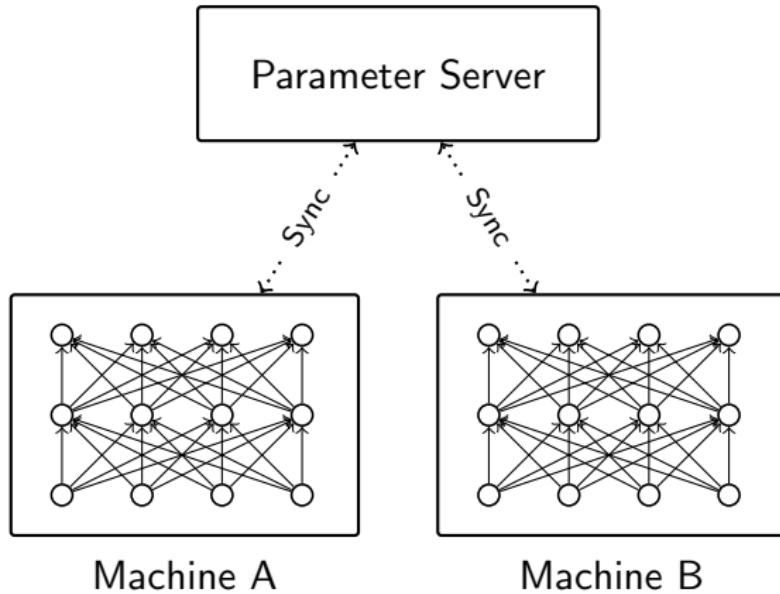


**Server<sub>0</sub>**



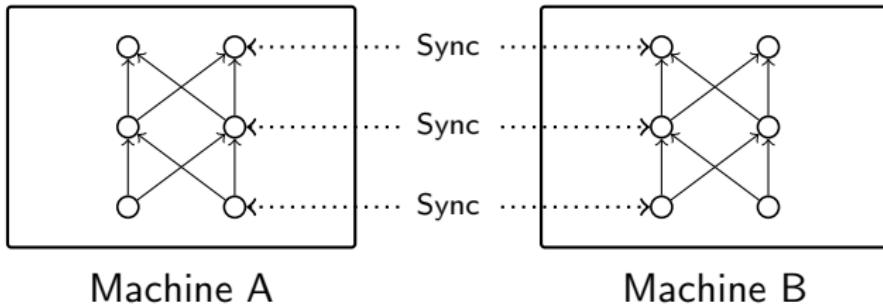
**Server<sub>1</sub>**

# Graph Layer Parallelism



## Data Parallelism

# Graph Layer Parallelism



## Model Parallelism

# Static Graphs

- ▶ Declarative programming model
- ▶ Static shapes and sizes
- ▶ Graphs defined once
- ▶ Can be compiled and optimized

```
Matrix<6, 9> a = ...;
Matrix<6, 9> b = ...;
Matrix<16, 6> c = ...;
scalar s = 7;

d := s * a + b;
e := matmul(c, d);

result := evaluate(e);
```

# Static Graphs

- ▶ Declarative programming model
- ▶ Static shapes and sizes
- ▶ Graphs defined once
- ▶ Can be compiled and optimized

```
x = if_clause(condition,  
                then_value,  
                else_value);  
  
while_loop([a, b],  
           condition_function,  
           body_function);
```

# Dynamic Graphs

- ▶ Imperative programming model
- ▶ Dynamic shapes and sizes
- ▶ Graphs defined by run

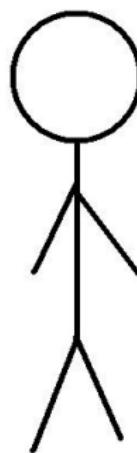
```
Matrix a = ...;
Matrix b = ...;
Matrix c = ...;
scalar s = 7;

if s > 0 {
    d := s * a + b;
} else {
    d := a - b;
}

while s > 0 {
    x = input();
    c = matmul(c, x);
}

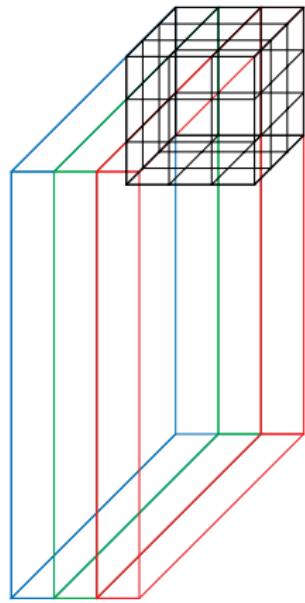
result := c;
```

# Op Layer

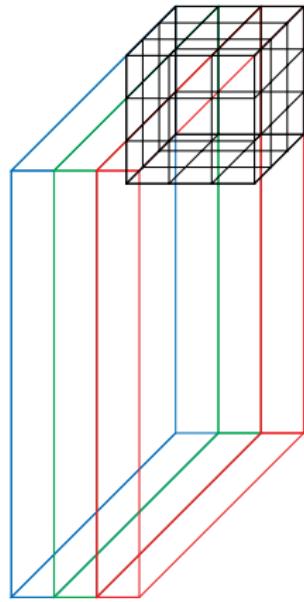


Bob

# im2col

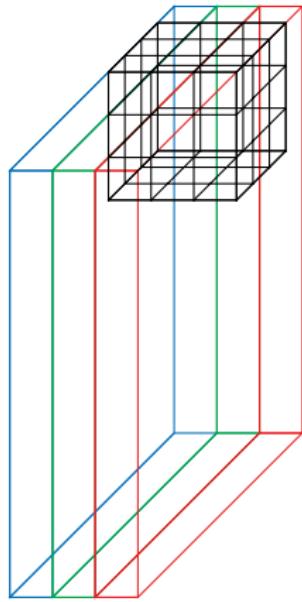


# im2col

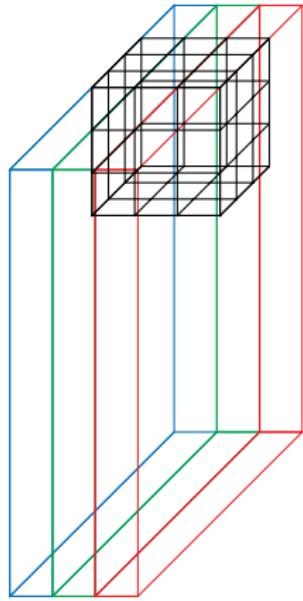


$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline \end{array} \quad \dots \quad \begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

# im2col

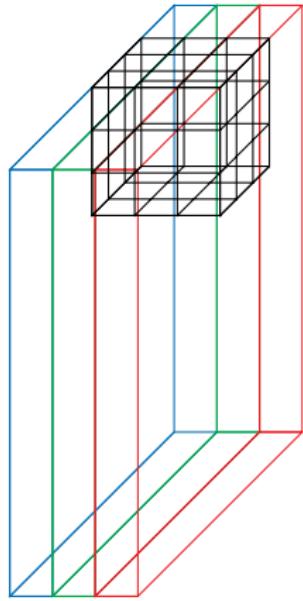

$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline x_0 & x_1 & x_2 \\ \hline \end{array} \quad \cdots \quad \begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

# im2col

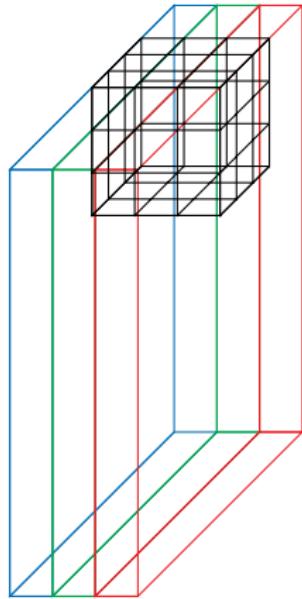


$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$

# im2col


$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline x_0 & x_1 & x_2 \\ \hline x_0 & x_1 & x_2 \\ \hline \vdots & \vdots & \vdots \\ \hline x_0 & x_1 & x_2 \\ \hline \end{array} \quad \cdots \quad \begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline x_{24} & x_{25} & x_{26} \\ \hline x_{24} & x_{25} & x_{26} \\ \hline \vdots & \vdots & \vdots \\ \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

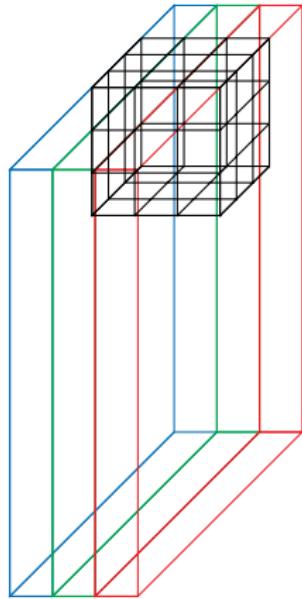
# im2col



$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$
⋮	⋮	⋮	⋮	⋮	⋮	⋮
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$

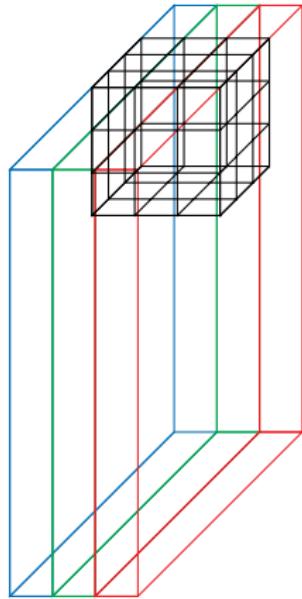
$k_0$
$k_1$
$k_2$
⋮
$k_{24}$
$k_{25}$
$k_{26}$

# im2col



$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$	$k_0$	$k_0$
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$	$k_1$	$k_1$
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$	$k_2$	$k_2$
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
$x_0$	$x_1$	$x_2$	⋮	$x_{24}$	$x_{25}$	$x_{26}$	$k_{24}$	$k_{24}$
							$k_{25}$	$k_{25}$
							$k_{26}$	$k_{26}$

# im2col



$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline \end{array}$$

$$\vdots \quad \vdots \quad \vdots$$

$$\begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline \end{array}$$

⋮

$$\begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

$$\vdots \quad \vdots \quad \vdots$$

$$\begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline k_0 & k_0 & k_0 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline k_1 & k_1 & k_1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline k_2 & k_2 & k_2 \\ \hline \end{array}$$

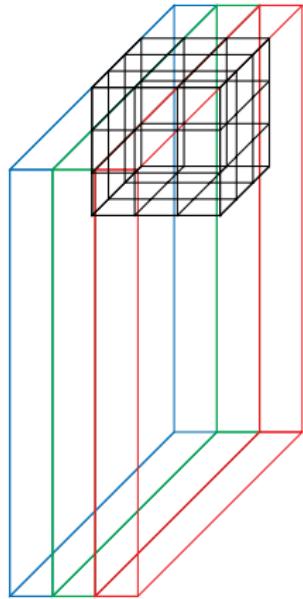
$$\vdots \quad \vdots \quad \vdots$$

$$\begin{array}{|c|c|c|} \hline k_{24} & k_{24} & k_{24} \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline k_{25} & k_{25} & k_{25} \\ \hline \end{array}$$

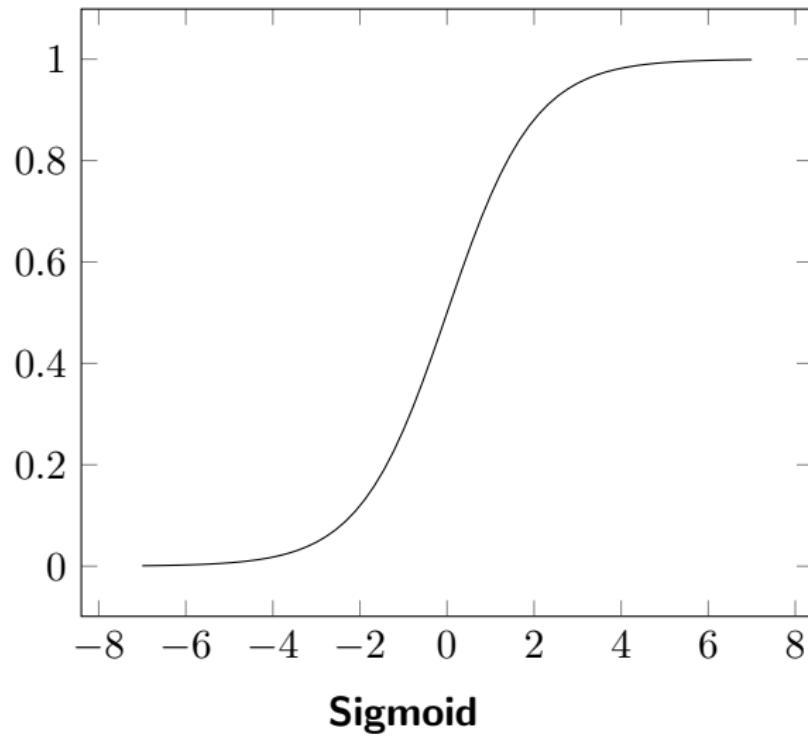
$$\begin{array}{|c|c|c|} \hline k_{26} & k_{26} & k_{26} \\ \hline \end{array}$$

# im2col



$$\begin{array}{c} \begin{array}{|c|c|c|} \hline x_0 & x_1 & x_2 \\ \hline x_0 & x_1 & x_2 \\ \hline x_0 & x_1 & x_2 \\ \hline \vdots & \vdots & \vdots \\ \hline x_0 & x_1 & x_2 \\ \hline \end{array} & \cdots & \begin{array}{|c|c|c|} \hline x_{24} & x_{25} & x_{26} \\ \hline x_{24} & x_{25} & x_{26} \\ \hline x_{24} & x_{25} & x_{26} \\ \hline \vdots & \vdots & \vdots \\ \hline x_{24} & x_{25} & x_{26} \\ \hline \end{array} & \times & \begin{array}{|c|c|c|} \hline k_0 & k_0 & k_0 \\ \hline k_1 & k_1 & k_1 \\ \hline k_2 & k_2 & k_2 \\ \hline \vdots & \vdots & \vdots \\ \hline k_{24} & k_{24} & k_{24} \\ \hline k_{25} & k_{25} & k_{25} \\ \hline k_{26} & k_{26} & k_{26} \\ \hline \end{array} \\ \end{array}$$

# Kernel Layer



# Kernel Layer

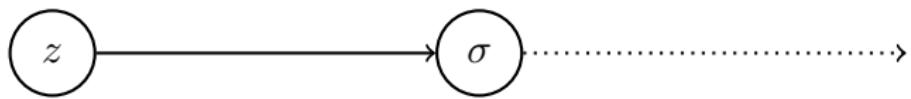
$$\sigma(z) = \frac{1}{1+e^{-z}}$$

**Forward Pass**

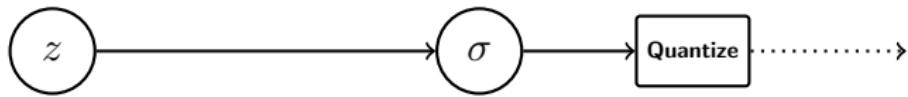
$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

**Backward Pass**

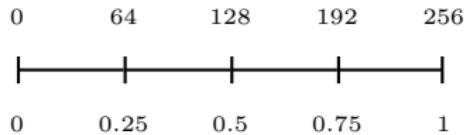
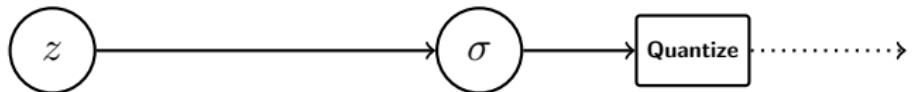
# Quantization



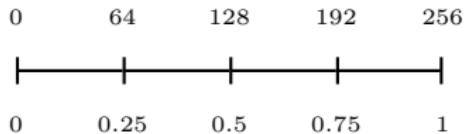
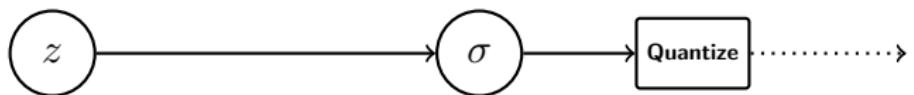
# Quantization



# Quantization



# Quantization



[github.com/google/gemmlowp](https://github.com/google/gemmlowp)

# Hardware Layer

# Hardware Layer



**CPU**

# Hardware Layer



**CPU**



**GPU**

# Hardware Layer



**CPU**

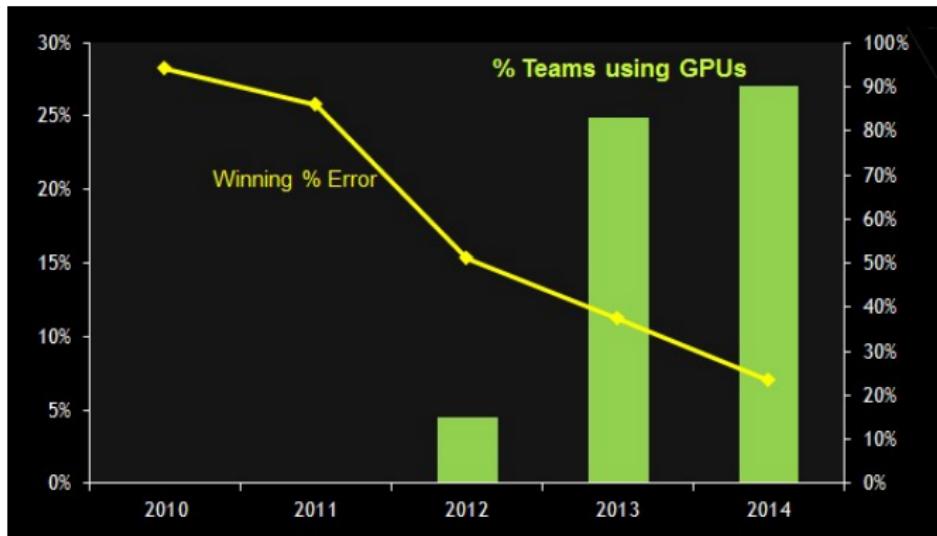


**GPU**



**ASIC**

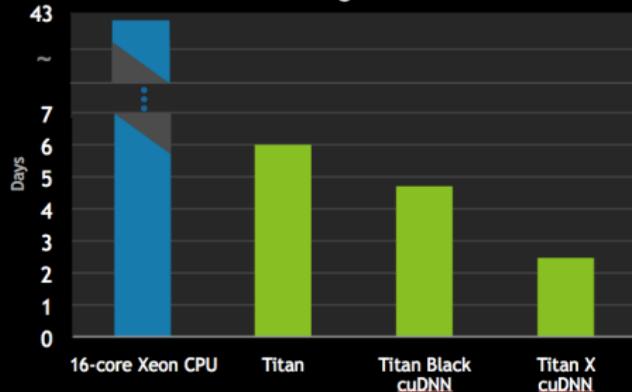
# GPUs



Teams using GPUs at ImageNet Competition

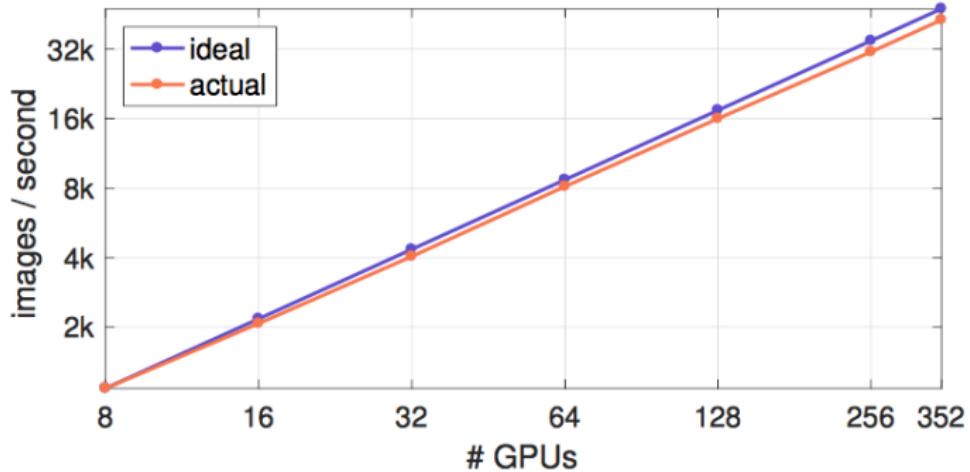
## TITAN X FOR DEEP LEARNING

Training AlexNet



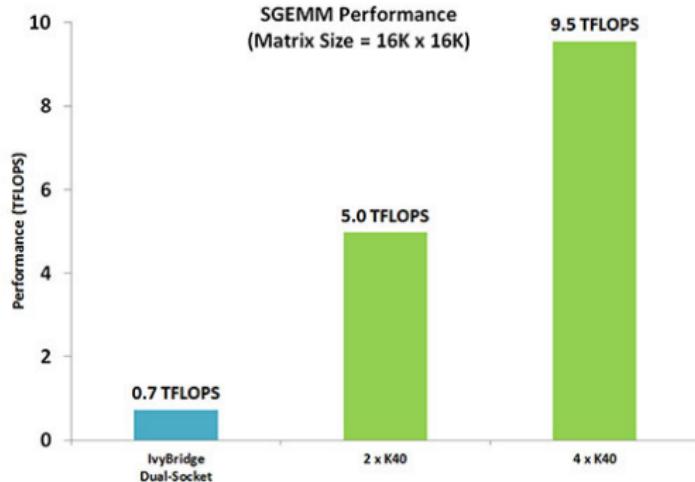
Training Time for ImageNet

# GPGUs



Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour  
Goyal et al. (2012)

# GPGPUs



Single Precision General Matrix/Matrix Multiply Performance

THIS IS YOUR MACHINE LEARNING SYSTEM?

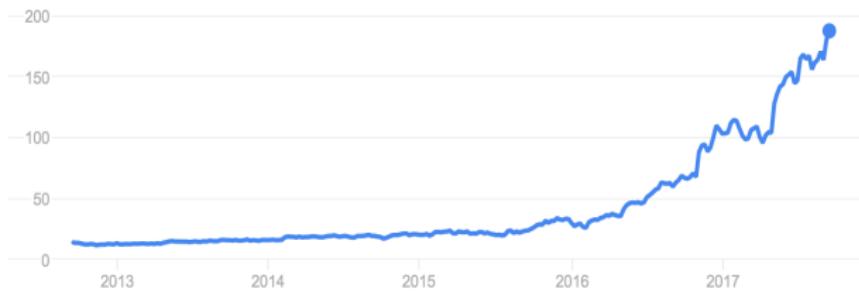
YUP! YOU POUR THE DATA INTO THIS BIG  
PILE OF LINEAR ALGEBRA, THEN COLLECT  
THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

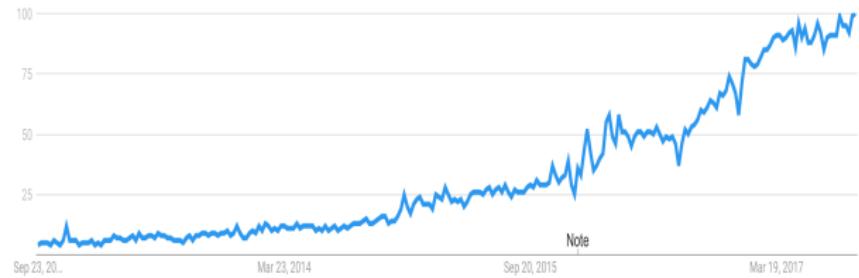
JUST STIR THE PILE UNTIL  
THEY START LOOKING RIGHT.



# Hardware: GPU



NVIDIA Stock



Deep Learning Trend

## Hardware: GPU

## Hardware: GPU



**Titan X**

# Hardware: GPU



**Titan X**



**Jetson TX2**

# Hardware: GPU



Titan X



Jetson TX2



DGX-1

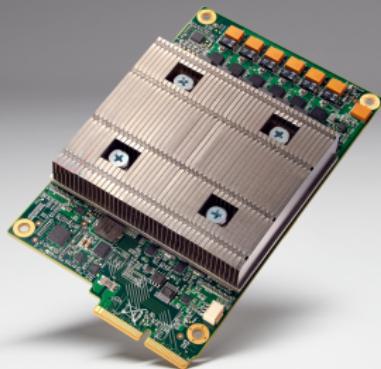
# Hardware: Big Basin



**Big Basin**

- ▶ 8 NVIDIA Tesla P100 GPUs
- ▶ NVLink (12x faster than PCIe)
- ▶ 16 GB RAM
- ▶ 10.6 TFLOPS
- ▶ Reduced Precision Computing

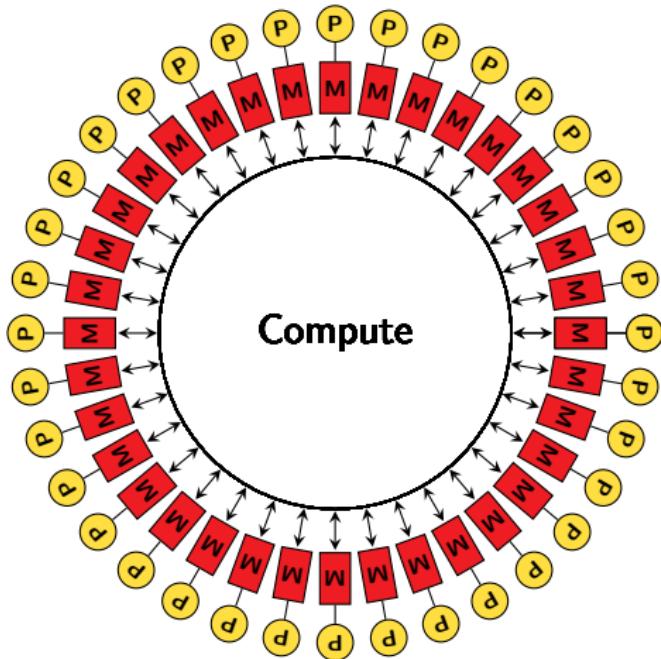
# Hardware: TPU



**TPU**

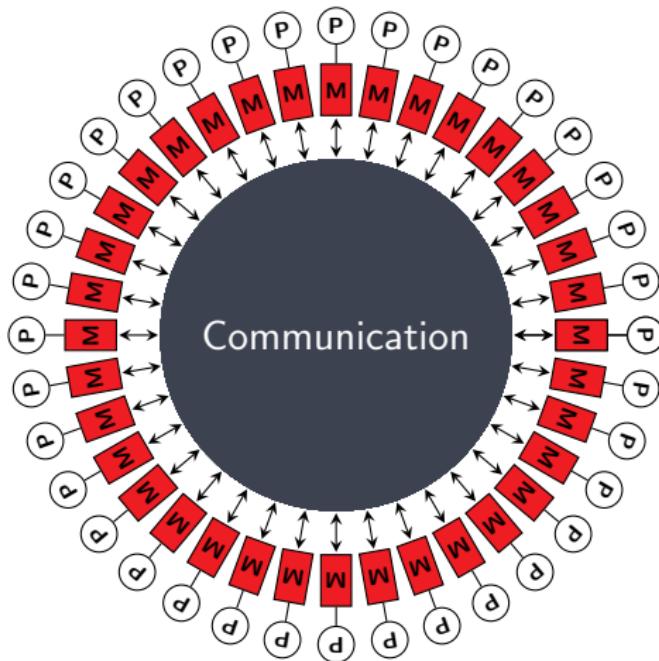
- ▶ *Coprocessor*
- ▶ 92 TOPS for 8-bit int
- ▶ 42 TFLOPS (TPU2)
- ▶ 24 MB Memory

## Hardware: IPU



**Bulk Synchronous Parallelism**

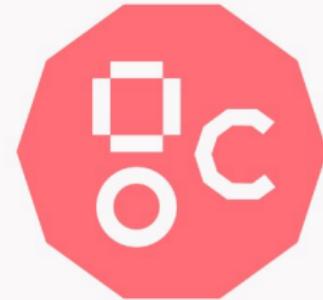
# Hardware: IPU



**Bulk Synchronous Parallelism**

## Hardware: IPU

- ▶ Startup from Bristol (UK)
- ▶ Graphcore Colossus
- ▶ TBR later this year
- ▶ 1000 Processors/Chip
- ▶ Mixed Precision Arithmetic
- ▶ No DRAM
- ▶ “Thrives on Sparsity”

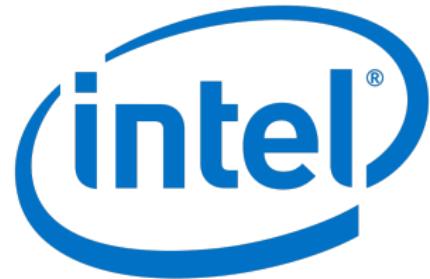


**Graphcore**

## Kernel Layer



cuDNN/cuBLAS



MKL

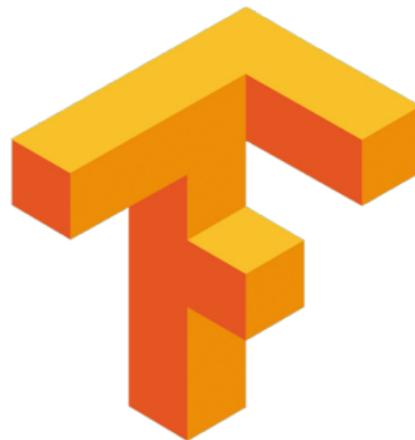
# **Demo:** Convolution w/ cuDNN

## Demo: TensorFlow Custom Op

# Graph Layer: TensorFlow

# Graph Layer: TensorFlow

- ▶ Google (2015)
- ▶ Static Graphs
- ▶ GPUs ✓
- ▶ Distributed ✓
- ▶ C++ API ✗
- ▶ Very “Deployable”



## Graph Layer: PyTorch/Caffe2



- ▶ Facebook (2016)
- ▶ Dynamic Graphs
- ▶ For Research
- ▶ GPUs ✓
- ▶ Distributed ✓
- ▶ Pure Python

- ▶ Facebook (2017)
- ▶ Static Graphs
- ▶ For Deployment
- ▶ GPUs ✓
- ▶ Distributed ✓
- ▶ C++ API ✓

## Graph Layer: MXNet

- ▶ Community, then Amazon
- ▶ Apache Incubator
- ▶ Highly Modular
- ▶ GPUs ✓
- ▶ Distributed ✓
- ▶ APIs for many languages  
(C++ ✓)



## Graph Layer: Others

Microsoft  
**CNTK**    **theano**



## Demo: MXNet Classifier

# Demo: Loading a TensorFlow graph

# How do I continue?

## Resources

- ▶ Machine Learning @ Coursera
- ▶ Deep Learning @ Coursera
- ▶ Deep Learning (Google) @ Udacity

## Resources

- ▶ Machine Learning @ Coursera
- ▶ Deep Learning @ Coursera
- ▶ Deep Learning (Google) @ Udacity
  
- ▶ [mxnet.incubator.apache.org/architecture](http://mxnet.incubator.apache.org/architecture)
- ▶ [github.com/dmlc/nnvm](https://github.com/dmlc/nnvm)
- ▶ Intro to Parallel Programming (NVIDIA) @ Udacity
- ▶ A Tour of TensorFlow, Goldsborough (2016)

## Resources

- ▶ Machine Learning @ Coursera
- ▶ Deep Learning @ Coursera
- ▶ Deep Learning (Google) @ Udacity
  
- ▶ [mxnet.incubator.apache.org/architecture](http://mxnet.incubator.apache.org/architecture)
- ▶ [github.com/dmlc/nnvm](https://github.com/dmlc/nnvm)
- ▶ Intro to Parallel Programming (NVIDIA) @ Udacity
- ▶ A Tour of TensorFlow, Goldsborough (2016)

[github.com/peter-can-talk/cppcon-2017](https://github.com/peter-can-talk/cppcon-2017)

# Q & A